



GameJam

Short Introduction on (Symbolic) AI Methods



Disclaimer

- This presentation focuses on the usage of symbolic AI techniques for level generation and action selection since they:
 - Often don't require that much computational power
 - No training and not that much parameter tuning
 - Allow good control over the desired outputs if the problems are well defined
- We don't talk about generative AI/Deep Learning techniques
 - Often way to compute intensive
 - Training not feasible in short projects
 - Hard to handle/control, especially if they need to fit into a larger concept



About us & Promotion

- Jakob Karalus - PhD Student at AI Institute - Human-in-the-Loop Reinforcement Learning & Explainability
- Conny Olz - Postdoc at AI Institute - Hierarchical Symbolic Planning

Self-Promotion: **Project “AI in Games”**

- Masters project at Ulm University
- Offered by the Institute of Artificial Intelligence
- Every semester, this summer term Tuesdays 14-16 pm
- Task: Develop a video game as a group using AI methods
 - Very similar to this GameJam

Self-Promotion: Project “AI in Games”



Self-Promotion: Project “AI in Games”





Outline

- Level Generation
 - Wave Function Collapse
 - Constraint Satisfaction Problem (CSP)
- Agent Movement
 - Pathfinding recap
 - Group movement with Boids
- Agent Action Selection
 - A* Search
 - Minimax

Generation with Wave Function Collapse



Wave Function Collapse - Idea

- Given:
 - An empty grid
 - A set of possible tiles
 - A set of constraints for the tiles.
- Example:
 - Water can be near water or beach (Blue)
 - Beach can be near beach, water or grass
 - Grass can be near beach, grass or hills
 - Hills can be near grass or hills
- Algorithms

Init grid with each value possible for each tile.

While not all tiles filled:

Select tile with lowest entropy and collapse it (if multiple, select randomly)

Propagate constraints




WFC - Example

Initialize empty grid



WFC - Example


Initialize every tile with all possible values (sometimes called waves)



WFC - Example

Select the lowest entropy cell -> All cell the same, so random.


Collapse the cell to one (still valid) value.



WFC - Example

Propagate constraints after collapsing.

Only water and beach can be near water.




WFC - Example

Propagate constraints after collapsing.

Since Hills can be only near water, a further set of tills can have some constraints collapsed.

Collapse constraints till no changes occur.




WFC - Example

Select the lowest entropy cell -> One of the 4 cells with only 2 constraints left.

Collapse the cell to one (still valid) value.

Only water/beach left, so 50/50 chance.




WFC - Example

Propagate constraints after collapsing.

Only water and beach can be near water.

Hills can only be near grass.

Collapse constraints till no changes occur.




WFC - Example

Select the lowest entropy cell.

Collapse the cell to one (still valid) value.

Only water/beach left, so 50/50 chance.

No propagation of constraints needed.




WFC - Example

Select the lowest entropy cell.


Collapse the cell to one (still valid) value.

Only water/beach left, so 50/50 chance.



WFC - Example

Update constraints.




WFC - Example


Select the lowest entropy cell.

Collapse the cell to one (still valid) value.

Only water/beach left, so 50/50 chance.



Update constraints




WFC - Example

Select the lowest entropy cell.


Collapse the cell to one (still valid) value.

Only water/beach left, so 50/50 chance.




WFC - Example

Update constraints



WFC - Example

Initialize empty grid



WFC - How to write constraints

- Writing all constraints can be tedious.
- Idea: Generate a example level which contains possibles allowed connections of tiles
- All NxN pattern in the output have to be in the input.
- Optional: The distribution of pattern should similar in the input and output pattern.
- Eg: <https://github.com/mxgmn/WaveFunctionCollapse>





WFC Online Demo

<https://oskarstalberg.com/game/wave/wave.html>



WFC Advantages

Advantages:

- WFC is quite easy to understand and simple to implement
- WFC can work with existing tiles (aka handplace a part of tiles)
- The constraints interface through example maps are quite convenient


Problems:

- Large states spaces
 - Multiple rooms in our dungeon crawler
- Constraints over a large amount of tiles
 - “There should exists a path between all rooms”
- Hard to specify example maps for large and different configuration
 - I want to have rooms with multiple biomes
- Abstract constraints
 - The player should not see the same monster type twice in a row

Constraints Satisfaction Problems

CSP

- Wave Function Collapse is just a special version of the “Constraint Satisfaction Problem”:
- CSP have
 - A set of states S - 2D Grid for the WFC, CSP could be any shape
 - A set of variables - Possible tile values in WFC, here similar
 - A set of constraints -
- One of the common textbook problems of CSPs is the Graph Coloring problems



Example 2: Sudoku

- Variables: X_{11}, \dots, X_{99} , one variable for every cell
- Domains: $\{1, \dots, 9\}$
- Constraints:
 - Row constraints: $X_{11} \neq X_{12}, \dots, X_{11} \neq X_{19}, \dots$
 - Column constraints: $X_{11} \neq X_{21}, \dots, X_{11} \neq X_{91}, \dots$
 - Block constraints: $X_{11} \neq X_{12}, \dots, X_{11} \neq X_{33}, \dots$

5	3			7				
6			1	9	5			
	9	8					6	
8				6				3
4			8		3			1
7				2				6
	6					2	8	
			4	1	9			5
				8			7	9

Example 3: Cryptarithmic Problem

- Variables: S, E, N, D, M, O, R, Y
- Domains: $\{0, \dots, 9\}$
- Constraints:
 - $S \neq E \neq N \neq D \neq M \neq O \neq R \neq Y$
 - $D + E = Y + 10 * C1$
 - $C1 + N + R = E + 10 * C10$
 - $C10 + E + O = N + 10 * C100$
 - $C100 + S + M = O + 10 * C1000$
 - $C1000 = M$
- (C1 ... C1000 are auxiliary variables representing the digit carried over)

$$\begin{array}{rccccccccc} & & & & S & & E & & N & & D & & \\ & & & & & & & & & & & & \\ + & & M & & O & & R & & E & & & & \\ \hline M & & O & & N & & E & & Y & & & & \end{array}$$



CSP Solving - Backtracking

Example:

1. A must be different from B.
 2. B must be greater than C.
 3. Possible values for A,B & C are {1,2,3}
-
- Start with an empty assignment: {}.
 - Start with variable A and value 1.
 - Check if the assignment satisfies all constraints:
 - Since A is not yet assigned, the first constraint is trivially satisfied.
 - There are no other variables assigned, so the second constraint is trivially satisfied.
 - Recursively assign values to the next variable (B).
 - Choose variable B and value 1.
 - Check if the assignment satisfies all constraints:
 - A is assigned 1, so the first constraint is violated.
 - There are no other variables assigned, so the second constraint is trivially satisfied.
 - Backtrack and try a different value for B.
 - Choose variable B and value 2.
 - Check if the assignment satisfies all constraints:
 - A is assigned 1, so the first constraint is satisfied.
 - There are no other variables assigned, so the second constraint is trivially satisfied.
 - Recursively assign values to the next variable (C).
 - Choose variable C and value 1.
 - Check if the assignment satisfies all constraints:
 - A is assigned 1, so the first constraint is satisfied.
 - B is assigned 2, so the second constraint is satisfied.
 - All variables are assigned values, so we have found a solution: {A: 1, B: 2, C: 1}.




WFC & CSP Summary

- CSP contain a number of more sophisticated solving technique, which allow for much more complicated constraint setup.
 - While techniques like backtracking, constraint propagation (forward checking) are simple, efficient implementation requires a bit of care
- In general WFC is better/faster in domains where a lot of solutions exists and require only local constraints. In settings where just a few solutions exists, classical CSP often outperform then.
- Classical CSP solver are a great complement to WFC style algorithms
 - E.g. Generate high-level graphs of room structure/biomes/enemies in a CSP. Low level tiling with WFC
- Software:
 - MiniZinc: <https://www.minizinc.org/> powerful CSP solver
 - Google OR Tools: <https://developers.google.com/optimization> collection of various optimization solvers
 - Kapsack problems - Fit the most item into a give space
 - Sheduling problems - Assign units to workspace to archive the highest utility
 - Traveling salesman - Shortest route to visit every location
 - ...

Agent Movement


Path finding

- Short recap A* and Dijkstra
- Probably already implemented in many game engines



Multiple Agents - Boids

- Individual path finding with possible overlapping routes can be computation expensive.
- Naive approach of just planning and then replanning leads to a lot of stuttering
- Boids:
 - Treat agents as a mass with the center moving into the right direction.





Boids - Pseudocode

Each entity is called a boid, with a position and a velocity (includes direction here)


Behaviour of boids is controlled by three factors:

- Separation
- Alignment
- Cohesion

```
def update_boids(boids):  
    """Update boids based on separation, alignment, and cohesion."""  
    for i, boid in enumerate(boids):  
        separation_force = separation(boids, i, SEPARATION_DISTANCE)  
        alignment_force = alignment(boids, i, ALIGNMENT_DISTANCE)  
        cohesion_force = cohesion(boids, i, COHESION_DISTANCE)  
  
        # Apply forces  
        boid.acceleration += separation_force * SEPARATION_WEIGHT  
        boid.acceleration += alignment_force * ALIGNMENT_WEIGHT  
        boid.acceleration += cohesion_force * COHESION_WEIGHT  
  
        # Update velocity and position  
        boid.velocity += boid.acceleration  
        boid.velocity.limit(MAX_SPEED)  
        boid.position += boid.velocity  
        boid.acceleration *= 0 # Reset acceleration
```

Boids - Pseudocode Separation


Each boid tries to keep a reasonable distance to all other boids in the same cluster.



```
def separation(boids, boid_index, separation_distance):
    """Calculate separation for a single boid."""
    steering = Vector(0, 0)
    count = 0
    for other_boid in boids:
        if other_boid is not boids[boid_index]:
            distance = distance_between(boids[boid_index].position, other_boid.position)
            if distance < separation_distance:
                diff = boids[boid_index].position - other_boid.position
                diff.normalize()
                diff /= distance
                steering += diff
                count += 1
    if count > 0:
        steering /= count
    if steering.length() > 0:
        steering.normalize()
        steering *= MAX_SPEED
        steering -= boids[boid_index].velocity
        steering.limit(MAX_FORCE)
    return steering
```

Boids - Pseudocode - Alignment


Each boid tries to align itself into the same direction than all other boids in his cluster.



```
def alignment(boids, boid_index, alignment_distance):  
    """Calculate alignment for a single boid."""  
    steering = Vector(0, 0)  
    count = 0  
    for other_boid in boids:  
        if other_boid is not boids[boid_index]:  
            distance = distance_between(boids[boid_index].position, other_boid.position)  
            if distance < alignment_distance:  
                steering += other_boid.velocity  
                count += 1  
    if count > 0:  
        steering /= count  
        steering.normalize()  
        steering *= MAX_SPEED  
        steering -= boids[boid_index].velocity  
        steering.limit(MAX_FORCE)  
    return steering
```

Boids - Pseudocode - Cohesion

Each boid tries to move into the direction of the mean positions of all other boids (in his cluster).




```
def cohesion(boids, boid_index, cohesion_distance):
    """Calculate cohesion for a single boid."""
    steering = Vector(0, 0)
    count = 0
    for other_boid in boids:
        if other_boid is not boids[boid_index]:
            distance = distance_between(boids[boid_index].position, other_boid.position)
            if distance < cohesion_distance:
                steering += other_boid.position
                count += 1
    if count > 0:
        steering /= count
        steering -= boids[boid_index].position
        steering.normalize()
        steering *= MAX_SPEED
        steering -= boids[boid_index].velocity
        steering.limit(MAX_FORCE)
    return steering
```

Agent Action Selection


Search-based Problems

- Represent the problem as tree/graph
 - States are nodes
 - Actions are edges
- Search through the tree for the best option
- Whole field of symbolic planning:
 - Idea: Just model the world, its states and actions and let the symbolic AI systems find the a action plan from the start to the goal
 - Requires a bit of reading, but can lead to really adaptive agents
 - See <https://www.fast-downward.org/>



A* - Search with a Heuristic


- Given a state/node, a heuristic estimates how many actions are still necessary to reach a goal state
 - E.g., count number of correct tiles
- A*: Expand node n with lowest f-value
 - $f(n) = g(n) + h(n)$
 - $g(n)$: costs to reach n
 - $h(n)$: heuristic value of n
- Guarantees to find optimal solution if heuristic fulfills certain properties (admissible)



Search - Multiple Players

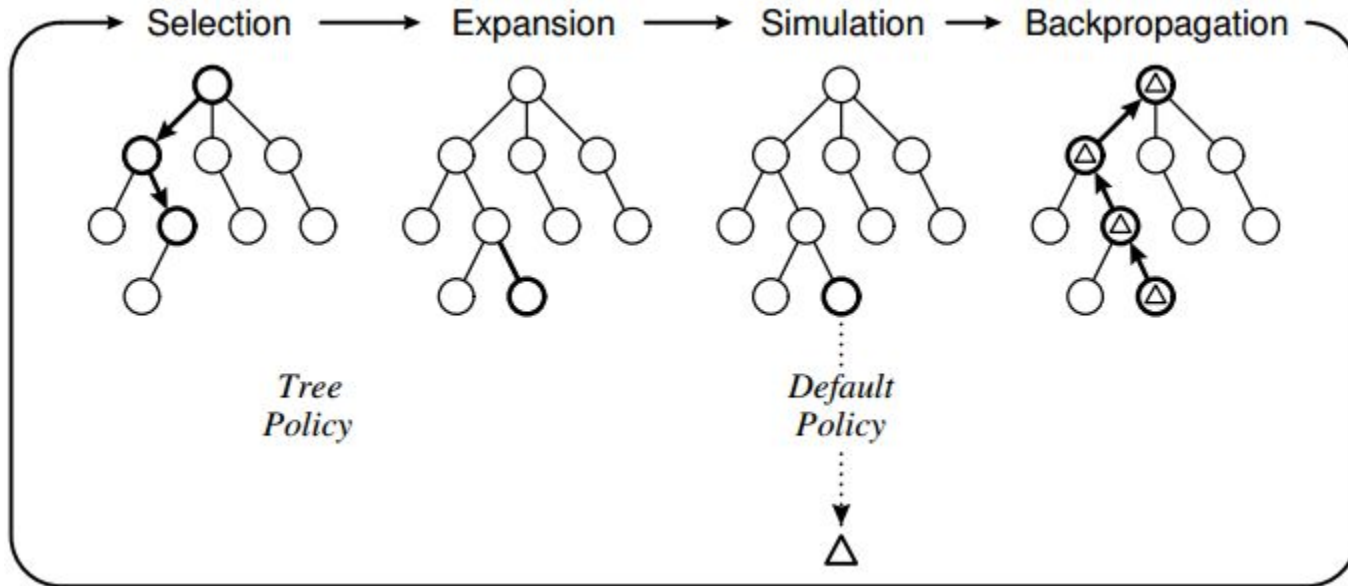
Minimax

- 2-Player games
- Assumptions: Fully observable, discrete, opponent plays optimal
- Question: What's the best move/strategy for winning the game?
- Solution:
 - Model the game as a search tree
 - Propagate results upwards
- Performance improvement: Alpha-Beta (pruning)
- Limitation: Large state spaces
 - > Monte Carlo Tree Search (next slides)



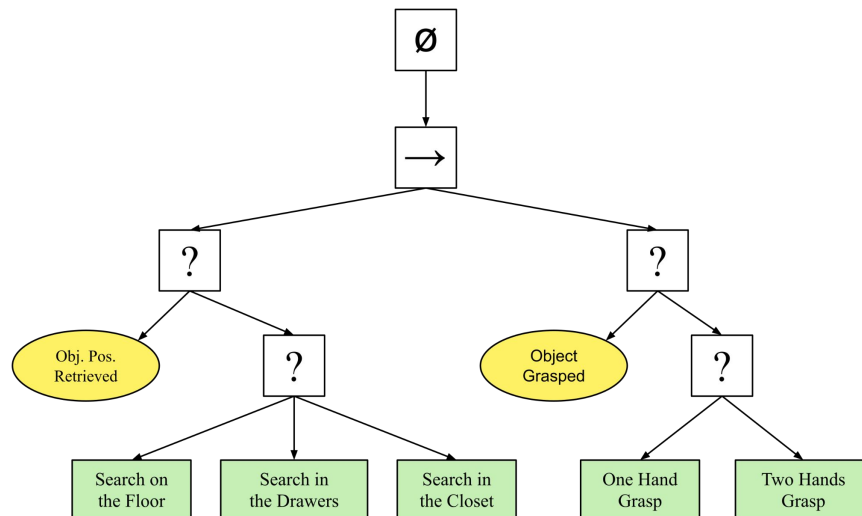
Very large search spaces

- If the search space get extremely large, sampling is needed.
- Probably waaaaay out of scope for this talk.



Behaviour Tree's

- For most games sophisticated search based action selection is probably overkill
- “Dumb Agents” doesn't necessary mean boring gameplay
 - Dark Souls have on the most scripted/simple reactive enemies...
- Direct often ascripting viable solution
- Behaviour Trees provide a bit more structure (especially in larger projects)
 - Combination between decisions trees and state machines



Realistic action selection - Thompson Sampling

- Always selecting the best option can lead to repetitive agents
- Random actions can lead to “stupid” agents
- Inspiration: Thompson Sampling:
 - Select an action proportional to how good it is
 - From the Multi-Armed-Bandit setting, but can be applied to any situation, either in search or direct action selection.
- How do we get to probabilities?
 - Softmax function
 - Can be adjusted through the temperature

```
import numpy as np
def softmax(logits, temperature=1):
    """
    Applies the softmax function to a list of logits.

    Args:
        logits: A list of floats representing unnormalized log probabilities.
        temperature: A float (default: 1) to control the "peakedness" of the distribution.
                    Higher temperature leads to a more uniform distribution.

    Returns:
        A list of floats representing the normalized probabilities.
    """
    e = np.exp(np.array(logits) / temperature)
    return e / np.sum(e)

def sample(probs):
    """
    Samples an index from a probability distribution.

    Args:
        probs: A list of floats representing probabilities.

    Returns:
        An integer representing the sampled index.
    """
    return np.random.choice(len(probs), p=probs)

# Example usage
logits = [1.0, 5.0, -0.1]
print(softmax(logits)) # Output: [0.01787917 0.97616938 0.00595146]
print(softmax(logits, temperature=3)) # Output: [0.18225863 0.69142873 0.12631264]
sample_index = sample(softmax(logits))
print(sample_index) # Output: An integer between 0 and 2 (inclusive)
```



Final notes

- Methods can be easily nested. Don't try to do everything with one method.
- Realistically for short project, hacking something together is often fine. For larger problems the usage of established tools can be highly beneficial.
-

Questions?
jakob.karalus@uni-ulm.de
